

5.ML_models - Hyperparameter

March 24, 2019

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import re
import time
import warnings
import sqlite3
from sqlalchemy import create_engine # database connection
import csv
import os
warnings.filterwarnings("ignore")
import datetime as dt
import numpy as np
from nltk.corpus import stopwords
from sklearn.decomposition import TruncatedSVD
from sklearn.preprocessing import normalize
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.manifold import TSNE
import seaborn as sns
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix
from sklearn.metrics.classification import accuracy_score, log_loss
from sklearn.feature_extraction.text import TfidfVectorizer
from collections import Counter
from scipy.sparse import hstack
from sklearn.multiclass import OneVsRestClassifier
from sklearn.svm import SVC
from sklearn.model_selection import StratifiedKFold
from collections import Counter, defaultdict
from sklearn.calibration import CalibratedClassifierCV
from sklearn.naive_bayes import MultinomialNB
from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV
import math
from sklearn.metrics import normalized_mutual_info_score
from sklearn.ensemble import RandomForestClassifier
```

```

from sklearn.model_selection import cross_val_score
from sklearn.linear_model import SGDClassifier
from mlxtend.classifier import StackingClassifier

from sklearn import model_selection
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import precision_recall_curve, auc, roc_curve

```

5. Assignments

1. Try out models (Logistic regression, Linear-SVM) with simple TF-IDF vectors instead of TD-IDF weighted word2Vec.
2. Hyperparameter tune XgBoost using RandomSearch to reduce the log-loss.

```

In [12]: if os.path.isfile('df_fe_without_preprocessing_train.csv'):
         dfppro = pd.read_csv("df_fe_without_preprocessing_train.csv",encoding='latin-1')

In [43]: train = pd.Series(df_vec['q1len'][0:40000].tolist() + df_vec['q2len'][0:40000].tolist()
         test = pd.Series(df_vec['q1len'][40001:60000].tolist() + df_vec['q2len'][40001:60000].tolist()

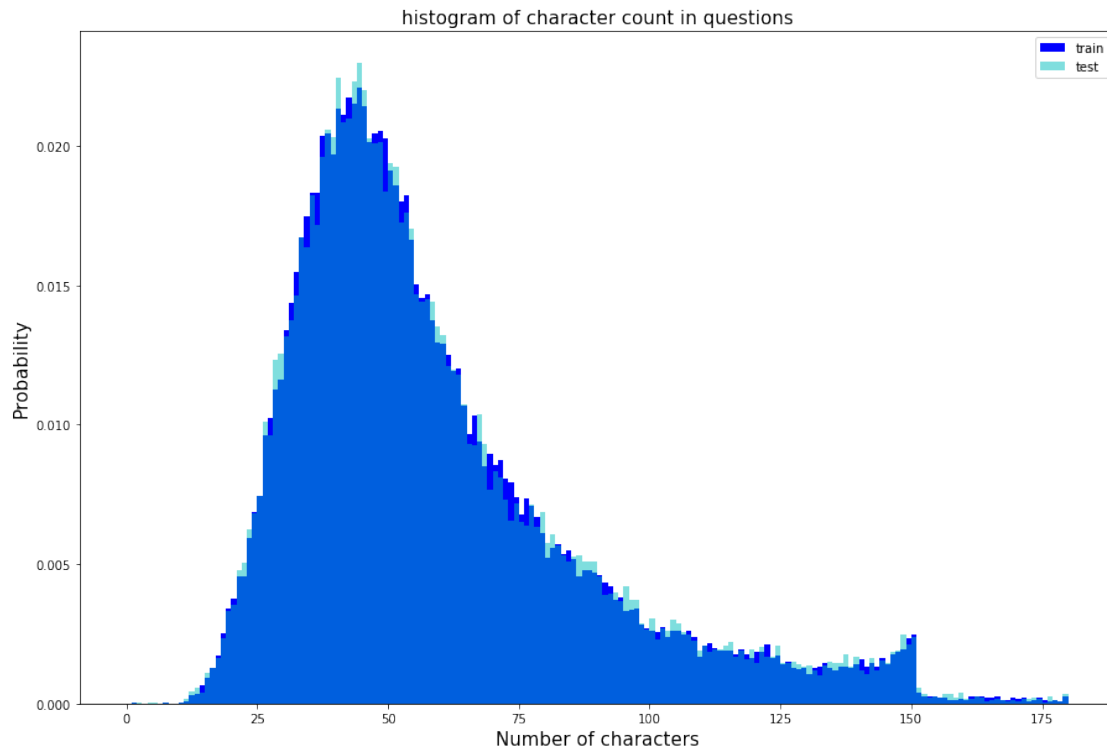
         plt.figure(figsize=(15, 10))
         plt.hist(train, bins=180, range=[0, 180], color='b', normed=True, label='train')
         plt.hist(test, bins=180, range=[0, 180], color='c', normed=True, alpha=0.5, label='test')
         plt.legend()
         plt.title('histogram of character count in questions', fontsize=15)
         plt.xlabel('Number of characters', fontsize=15)
         plt.ylabel('Probability', fontsize=15)
         plt.show()

         print ('The mean of length of questions in train is: {} and test is {}'.format(np.mean(train), np.mean(test)))

```

c:\users\dell\appdata\local\programs\python\python36\lib\site-packages\matplotlib\axes_axes.py:160: The 'normed' kwarg was deprecated in Matplotlib 2.1 and will be removed in 3.1. Use 'density' instead.
alternative="density", removal="3.1")

c:\users\dell\appdata\local\programs\python\python36\lib\site-packages\matplotlib\axes_axes.py:160: The 'normed' kwarg was deprecated in Matplotlib 2.1 and will be removed in 3.1. Use 'density' instead.
alternative="density", removal="3.1")



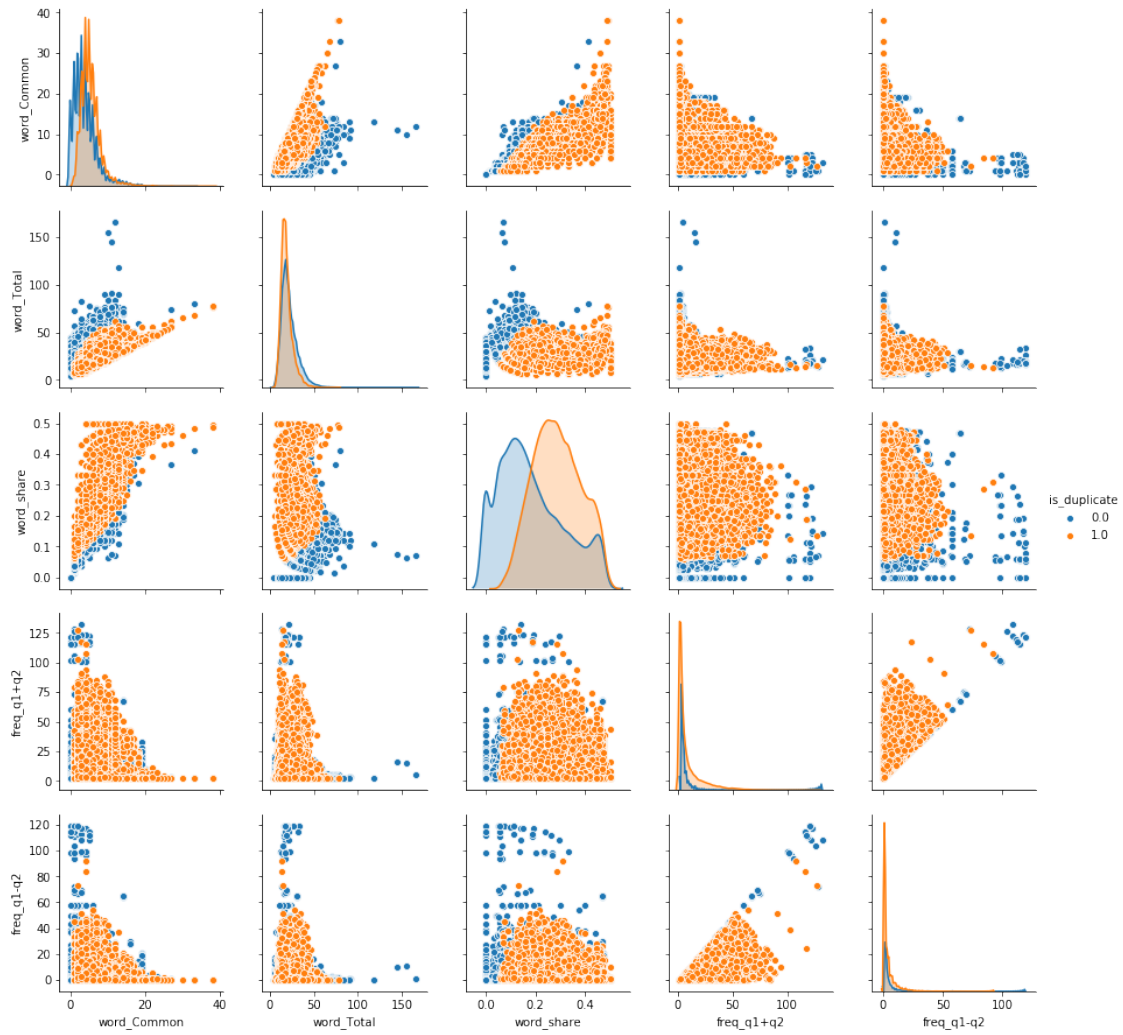
The mean of length of questions in train is: 59.6954 and test is 59.826266313315664

We can see that both train and test follow the same distribution and mean is around 60

0.1 Pair plots of basic features

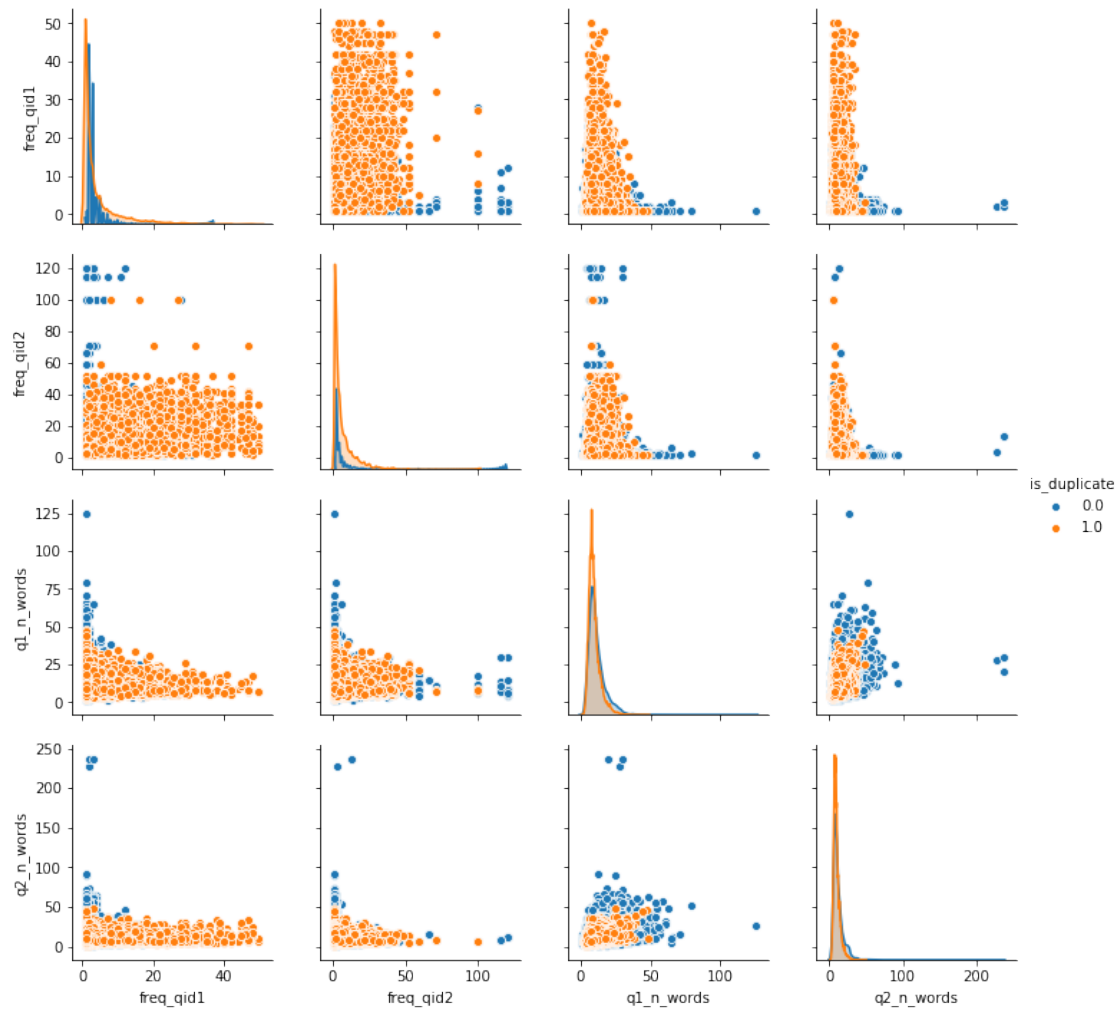
```
In [55]: #n = df.shape[0]
sns.pairplot((dfppro[['word_Common', 'word_Total', 'word_share', 'freq_q1+q2', 'freq_q1-q2']])
```

```
Out[55]: <seaborn.axisgrid.PairGrid at 0x29dc38ce860>
```



```
In [57]: sns.pairplot((dfppro[['freq_qid1','freq_qid2','q1_n_words','q2_n_words']]).join(dfnlp
```

```
Out[57]: <seaborn.axisgrid.PairGrid at 0x29df42d7ba8>
```



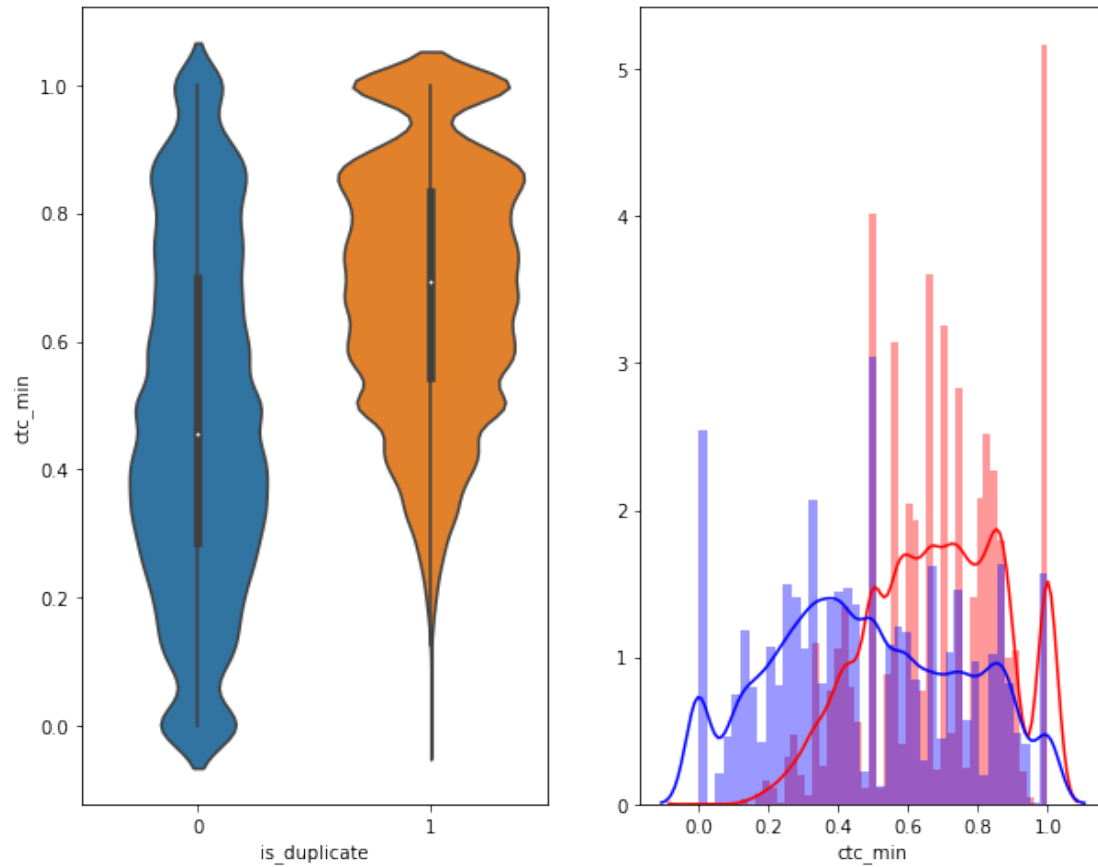
The features are important in classifying the points.

```
In [60]: dfc=dfppro[['word_Common','word_Total','word_share','freq_q1+q2','freq_q1-q2']].join(

In [65]: # Distribution of the ctc_min
plt.figure(figsize=(10, 8))

plt.subplot(1,2,1)
sns.violinplot(x = 'is_duplicate', y = 'ctc_min', data = dfnlp[['is_duplicate','ctc_min']])

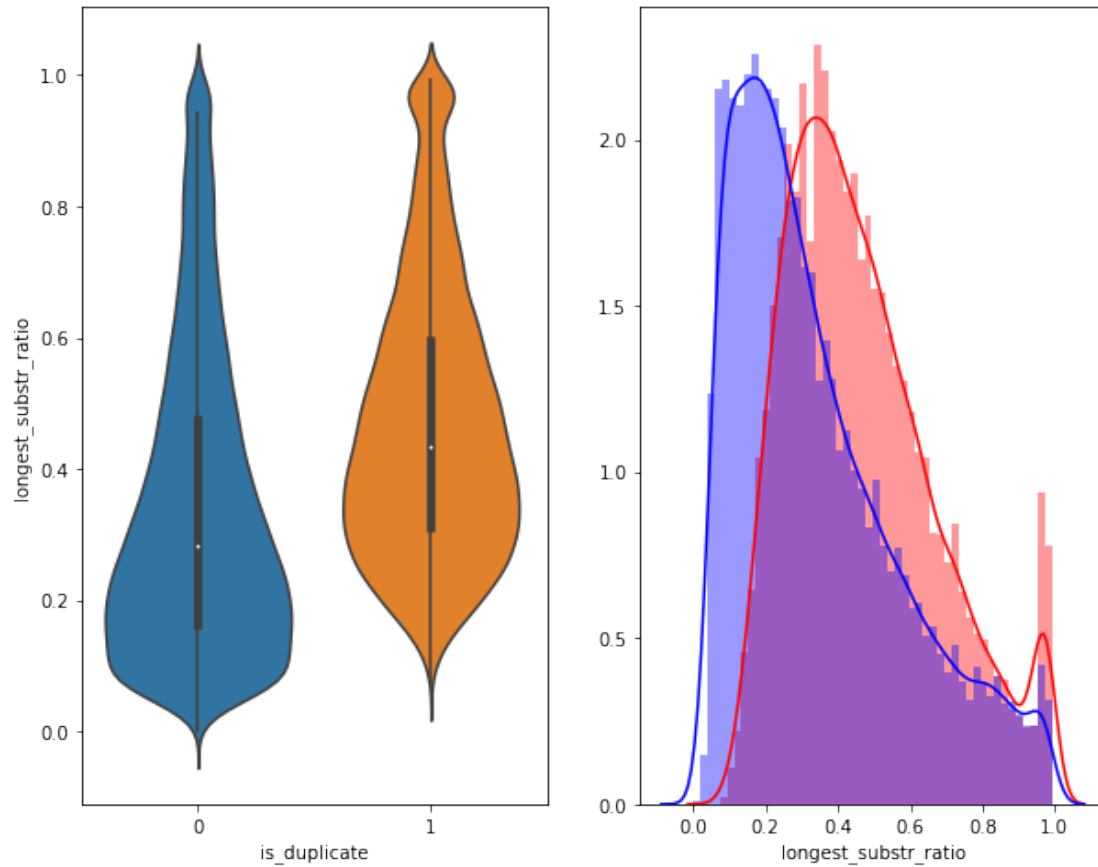
plt.subplot(1,2,2)
sns.distplot(dfnlp[dfnlp['is_duplicate'] == 1.0]['ctc_min'][0:], label = "1", color = "orange")
sns.distplot(dfnlp[dfnlp['is_duplicate'] == 0.0]['ctc_min'][0:], label = "0", color = "blue")
plt.show()
```



```
In [69]: # Distribution of the longest_substr_ratio
plt.figure(figsize=(10, 8))

plt.subplot(1,2,1)
sns.violinplot(x = 'is_duplicate', y = 'longest_substr_ratio', data = dfnlp[['is_dupli

plt.subplot(1,2,2)
sns.distplot(dfnlp[dfnlp['is_duplicate'] == 1.0]['longest_substr_ratio'][0:], label =
sns.distplot(dfnlp[dfnlp['is_duplicate'] == 0.0]['longest_substr_ratio'][0:], label =
plt.show()
```



In [70]: *# Distribution of the fuzz_partial_ratio*

```
plt.figure(figsize=(10, 8))
```

```
plt.subplot(1,2,1)
```

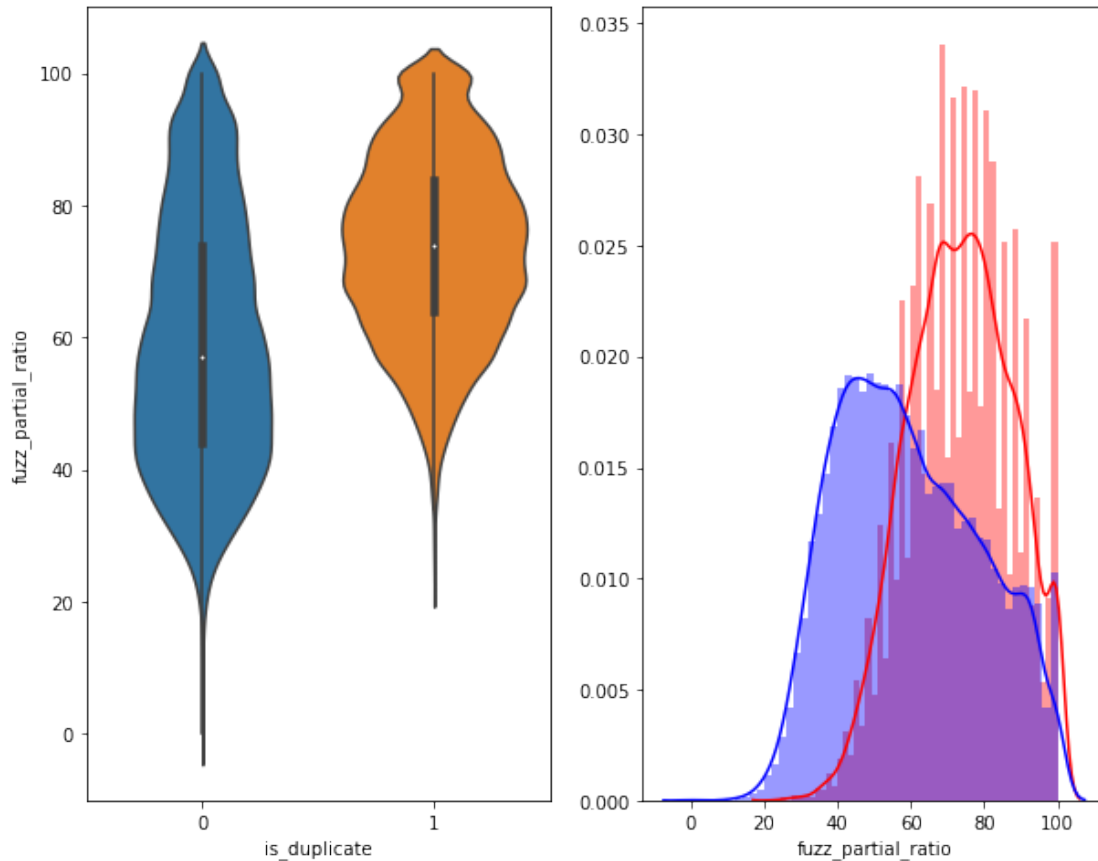
```
sns.violinplot(x = 'is_duplicate', y = 'fuzz_partial_ratio', data = dfnlp[['is_duplicate', 'fuzz_partial_ratio']])
```

```
plt.subplot(1,2,2)
```

```
sns.distplot(dfnlp[dfnlp['is_duplicate'] == 1.0]['fuzz_partial_ratio'][0:], label = 'is_duplicate = 1.0')
```

```
sns.distplot(dfnlp[dfnlp['is_duplicate'] == 0.0]['fuzz_partial_ratio'][0:], label = 'is_duplicate = 0.0')
```

```
plt.show()
```



0.2 Wordcloud of highest tfidf feature corresponding to y=0/1

```
In [83]: idx=dfnlp.index[dfnlp['is_duplicate'] == 1]
         positive=dfnlp.loc[idx, ['question1', 'question2']]
         idz=dfnlp.index[dfnlp['is_duplicate'] == 0]
         negative=dfnlp.loc[idz, ['question1', 'question2']]

In [95]: #vectorizing the questions corresponding to y=1
         tfidf_vect = TfidfVectorizer(ngram_range=(1,2), min_df=10)
         vec_pos=tfidf_vect.fit_transform(positive['question1']+positive['question2'])

In [97]: #getting the features with top tfidf values
         features=pd.DataFrame(tfidf_vect.get_feature_names())
         mean_tf=np.mean(vec_pos,axis=0)
         x=np.array(mean_tf)[0].tolist()
         # taking top 20 features
         important_feat=[]
         important_feat=np.argsort((x))[::-1]
         important_feat=important_feat[:100]
         important_feat
```



```

In [105]: #getting the features with top tfidf values
features=pd.DataFrame(tf_idf_vect.get_feature_names())
mean_tf=np.mean(vec_pos,axis=0)
x=np.array(mean_tf)[0].tolist()
# taking top 20 features
important_feat=[]
important_feat=np.argsort((x))[:, :-1]
important_feat=important_feat[:100]
important_feat
imp_feat=[]
for index in important_feat:
    imp_feat.append(features.iloc[index])

x=np.array(imp_feat)
feature=[]
for i in x:
    for j in i:
        feature.append(j)

In [106]: from wordcloud import WordCloud
# Ploting word cloud
# Lets first convert the 'result' dictionary to 'list of tuples'
#tup = dict(result.items())
#Initializing WordCloud using frequencies of tags.
wordcloud = WordCloud(    background_color='black',
                        width=1600,
                        height=800,
                        ).generate(str(feature))

fig = plt.figure(figsize=(30,20))
plt.imshow(wordcloud)
plt.axis('off')
plt.tight_layout(pad=0)
fig.savefig("tag.png")
plt.show()

```



```
fuzz_ratio          False
fuzz_partial_ratio  False
longest_substr_ratio False
dtype: bool
```

```
In [8]: #splitting manually into train and test
dfnlp=dfnlp[0:60000]
dfnlp.shape
dfnlp_train=dfnlp[0:40000]
dfnlp_test=dfnlp[40001:60000]
```

```
In [16]: print(dfnlp_train.shape)
         dfnlp_test.shape
```

```
(40000, 21)
```

```
Out[16]: (19999, 21)
```

```
In [17]: # tfidf vectorization with n_gram=2
tf_idf_vect = TfidfVectorizer(ngram_range=(1,2), min_df=10)
train_vec_1=tf_idf_vect.fit_transform(dfnlp_train['question1'])
test_vec_1=tf_idf_vect.transform(dfnlp_test['question1'])
train_vec_2=tf_idf_vect.fit_transform(dfnlp_train['question2'])
test_vec_2=tf_idf_vect.transform(dfnlp_test['question2'])
```

```
In [18]: if os.path.isfile('df_fe_without_preprocessing_train.csv'):
         dfppro = pd.read_csv("df_fe_without_preprocessing_train.csv",encoding='latin-1')
```

```
In [19]: #stacking the train vectors of question 1 and two
         from scipy import sparse
         train_q=sparse.hstack([train_vec_1,train_vec_2])
         #stacking the test vectors of question 1 and two
         test_q=sparse.hstack([test_vec_1,test_vec_2])
         train_q.shape
```

```
Out[19]: (40000, 14983)
```

```
In [20]: #getting rest of advanced features
         df_vec=dfnlp[['cwc_min','cwc_max','csc_min','csc_max','ctc_min','ctc_max','last_word_
```

```
In [21]: #y labels
         train_y=dfnlp['is_duplicate'][0:40000]
         test_y=dfnlp['is_duplicate'][40001:60000]
```

```
In [22]: #stacking advanced features and tfidf vectors
         train_vecs=sparse.hstack([train_q,df_vec[0:40000]])
         test_vecs=sparse.hstack([test_q,df_vec[40001:60000]])
```

```
In [23]: train_vecs.shape
```

```
Out[23]: (40000, 15009)
```

0.4 Logistic Regression with hyperparameter tuning

```
In [85]: from tqdm.auto import tqdm
alpha = [10 ** x for x in range(-5, 2)] # hyperparam for SGD classifier.

log_error_array=[]
for i in tqdm(alpha):
    clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=42)
    clf.fit(train_vecs, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_vecs, train_y)
    predict_y = sig_clf.predict_proba(test_vecs)
    log_error_array.append(log_loss(test_y, predict_y, labels=clf.classes_, eps=1e-15))
    print('For values of alpha = ', i, "The log loss is:", log_loss(test_y, predict_y,

fig, ax = plt.subplots()
ax.plot(alpha, log_error_array, c='g')
for i, txt in enumerate(np.round(log_error_array, 3)):
    ax.annotate((alpha[i], np.round(txt, 3)), (alpha[i], log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

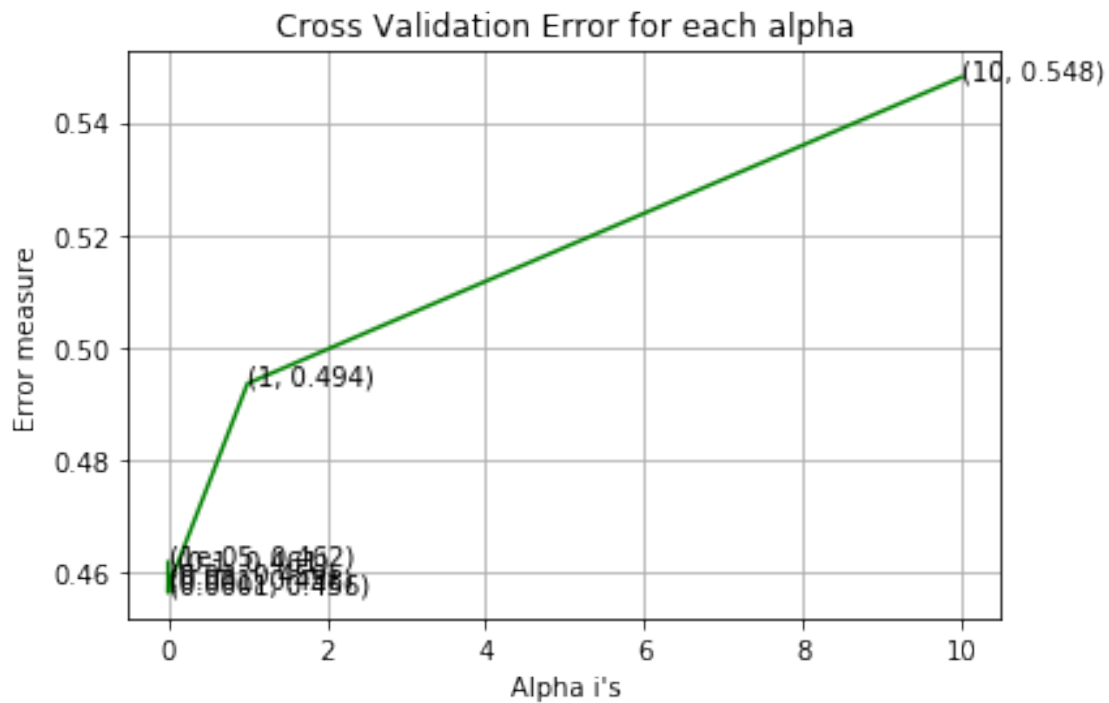
best_alpha = np.argmin(log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
clf.fit(train_vecs, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_vecs, train_y)

predict_y = sig_clf.predict_proba(train_vecs)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log_loss(
predict_y = sig_clf.predict_proba(test_vecs)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_loss(
predicted_y = np.argmax(predict_y, axis=1)
print("Total number of data points :", len(predicted_y))
plot_confusion_matrix(test_y, predicted_y)

HBox(children=(IntProgress(value=0, max=7), HTML(value='')))
```

For values of alpha = 1e-05 The log loss is: 0.46184114657478476
For values of alpha = 0.0001 The log loss is: 0.4564201176731175
For values of alpha = 0.001 The log loss is: 0.45756864449921325
For values of alpha = 0.01 The log loss is: 0.4585507362678816
For values of alpha = 0.1 The log loss is: 0.460880908451455
For values of alpha = 1 The log loss is: 0.4935858896035705

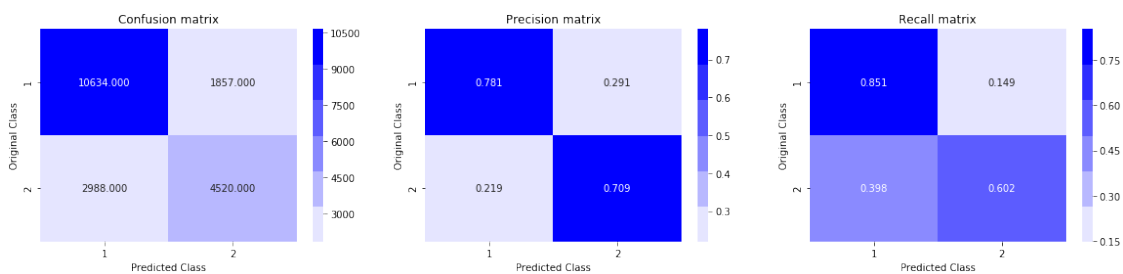
For values of alpha = 10 The log loss is: 0.5480599551149273



For values of best alpha = 0.0001 The train log loss is: 0.45278879999465105

For values of best alpha = 0.0001 The test log loss is: 0.4564201176731175

Total number of data points : 19999



0.5 Linear SVM with hyperparameter tuning

In [86]: alpha = [10 ** x for x in range(-5, 2)] # hyperparam for SGD classifier.

--

```

log_error_array=[]
for i in tqdm(alpha):
    clf = SGDClassifier(alpha=i, penalty='l1', loss='hinge', random_state=42)
    clf.fit(train_vecs, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_vecs, train_y)
    predict_y = sig_clf.predict_proba(test_vecs)
    log_error_array.append(log_loss(test_y, predict_y, labels=clf.classes_, eps=1e-15))
    print('For values of alpha = ', i, "The log loss is:", log_loss(test_y, predict_y,

fig, ax = plt.subplots()
ax.plot(alpha, log_error_array, c='g')
for i, txt in enumerate(np.round(log_error_array, 3)):
    ax.annotate((alpha[i], np.round(txt, 3)), (alpha[i], log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

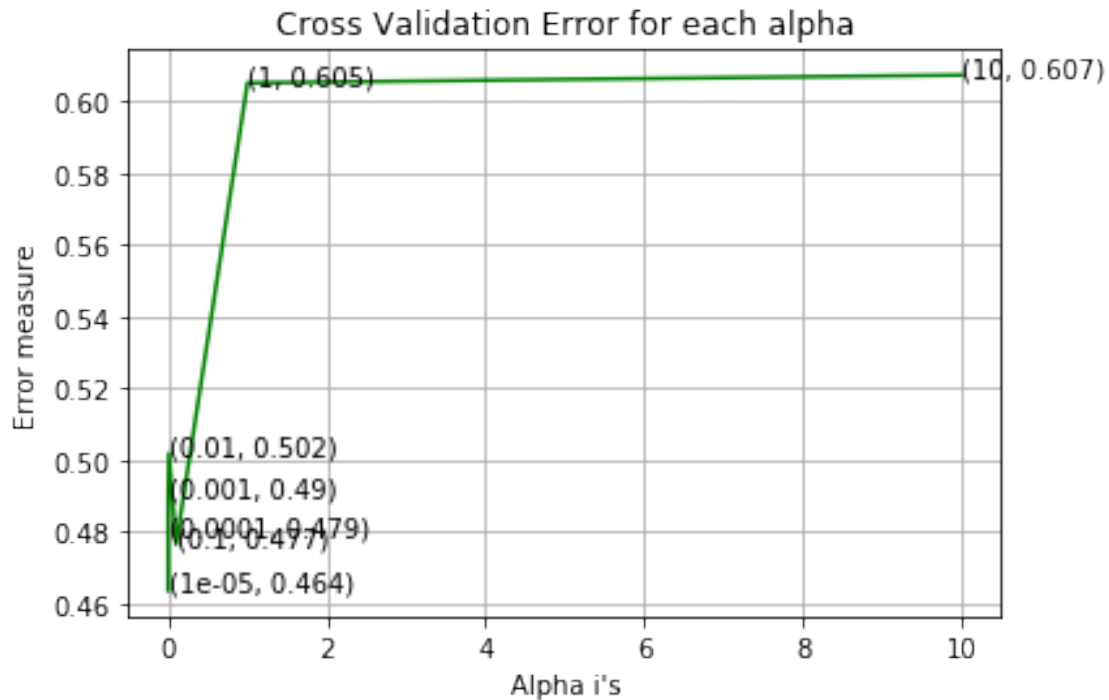
best_alpha = np.argmin(log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l1', loss='hinge', random_state=42)
clf.fit(train_vecs, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_vecs, train_y)

predict_y = sig_clf.predict_proba(train_vecs)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log_loss(train_y, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(test_vecs)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_loss(test_y, predict_y, labels=clf.classes_, eps=1e-15))
predicted_y = np.argmax(predict_y, axis=1)
print("Total number of data points :", len(predicted_y))
plot_confusion_matrix(test_y, predicted_y)

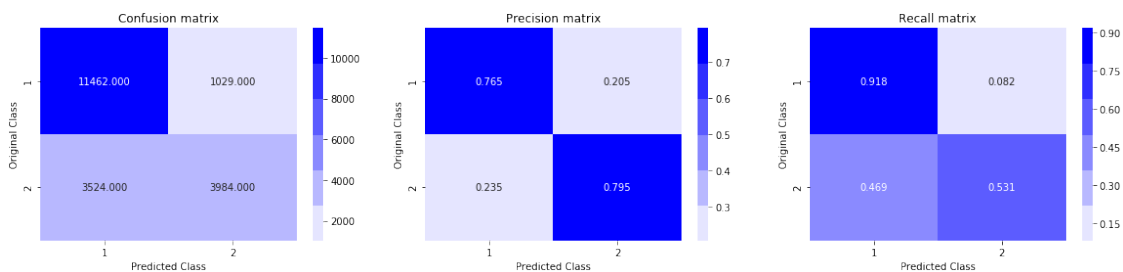
HBox(children=(IntProgress(value=0, max=7), HTML(value=''))))

For values of alpha = 1e-05 The log loss is: 0.4636498119106351
For values of alpha = 0.0001 The log loss is: 0.4789319096072303
For values of alpha = 0.001 The log loss is: 0.48969325832950694
For values of alpha = 0.01 The log loss is: 0.5018947372269298
For values of alpha = 0.1 The log loss is: 0.4765511555383432
For values of alpha = 1 The log loss is: 0.6052293455226234
For values of alpha = 10 The log loss is: 0.6074137784736146

```



For values of best alpha = 1e-05 The train log loss is: 0.45968365849046866
 For values of best alpha = 1e-05 The test log loss is: 0.4636498119106351
 Total number of data points : 19999



0.6 XGBOOST

```
In [151]: import xgboost as xgb
          # with glove vectorization
          from scipy.stats import randint as sp_randint
          from sklearn.model_selection import RandomizedSearchCV
          g=sp_randint(1, 50)
          # using randomsearch for hyperparameter tuning
```



```

xg=xgb.XGBClassifier(nthread=-1, early_stopping_rounds=20, verbose_eval=10)
params = {}
params['n_estimators']=10,100]
params['max_depth'] = sp_randint(1, 50)

random_search = RandomizedSearchCV(xg,cv=5, param_distributions=params,\
                                   n_jobs=-1,verbose=10 ,scoring='neg_log_loss')
random_search.fit(X_train,y_train)

```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

```

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done   2 tasks      | elapsed: 43.4min
[Parallel(n_jobs=-1)]: Done   9 tasks      | elapsed: 67.9min
[Parallel(n_jobs=-1)]: Done  16 tasks      | elapsed: 92.7min
[Parallel(n_jobs=-1)]: Done  25 tasks      | elapsed: 122.2min
[Parallel(n_jobs=-1)]: Done  34 tasks      | elapsed: 142.2min
[Parallel(n_jobs=-1)]: Done  41 out of  50 | elapsed: 148.2min remaining: 32.5min
[Parallel(n_jobs=-1)]: Done  47 out of  50 | elapsed: 164.2min remaining: 10.5min
[Parallel(n_jobs=-1)]: Done  50 out of  50 | elapsed: 166.7min finished

```

```

Out[151]: 'estimators=[]\n    depth=[]\n    key_scores={} \n\n    #storing the results\n    all

```

```

In [174]: random_search.best_params_

```

```

Out[174]: {'max_depth': 7, 'n_estimators': 100}

```

```

In [153]: #building a model with depth=7 and n_estimators=100

```

```

import xgboost as xgb
params = {}
params['objective'] = 'binary:logistic'
params['eval_metric'] = 'logloss'
params['eta'] = 0.02
params['max_depth'] = 7
params['n_estimators']=100

d_train = xgb.DMatrix(X_train, label=y_train)
d_test = xgb.DMatrix(X_test, label=y_test)

watchlist = [(d_train, 'train'), (d_test, 'valid')]

bst = xgb.train(params, d_train, 400, watchlist, early_stopping_rounds=20, verbose_e

xgdmatrix = xgb.DMatrix(X_train, label=y_train)
predict_y = bst.predict(d_test)
print("The test log loss is:",log_loss(y_test, predict_y, labels=clf.classes_, eps=1

```

```
[12:23:12] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[0]      train-logloss:0.683169      valid-logloss:0.683406
Multiple eval metrics have been passed: 'valid-logloss' will be used for early stopping.
```

Will train until valid-logloss hasn't improved in 20 rounds.

```
[12:23:14] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:23:16] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:23:17] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:23:19] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:23:20] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:23:22] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:23:23] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:23:25] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:23:26] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:23:28] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[10]      train-logloss:0.600874      valid-logloss:0.603531
[12:23:29] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:23:31] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:23:32] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:23:34] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:23:36] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:23:37] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:23:39] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:23:41] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:23:43] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:23:45] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[20]      train-logloss:0.541371      valid-logloss:0.546414
[12:23:47] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:23:49] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:23:50] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:23:52] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:23:54] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:23:56] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:23:58] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:23:59] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:24:01] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:24:03] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[30]      train-logloss:0.497047      valid-logloss:0.504672
[12:24:05] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:24:07] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:24:08] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:24:10] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:24:12] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:24:14] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:24:16] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:24:18] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:24:19] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:24:21] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
```


[illegible]

[illegible]


```

[12:35:54] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[390]      train-logloss:0.239867      valid-logloss:0.335152
[12:35:56] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:35:58] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:36:00] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:36:02] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:36:04] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:36:06] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:36:07] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:36:09] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[12:36:11] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[399]      train-logloss:0.237149      valid-logloss:0.334897
The test log loss is: 0.33489747420122634

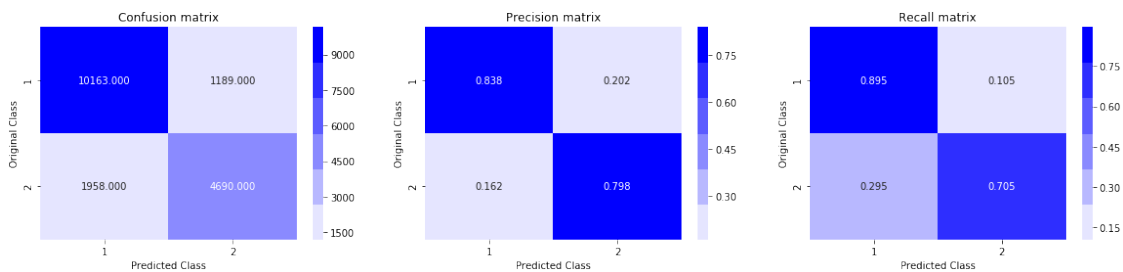
```

```

In [173]: predicted_y =np.array(predict_y>0.5,dtype=int)
          print("Total number of data points :", len(predicted_y))
          plot_confusion_matrix(y_test, predicted_y)

```

Total number of data points : 18000



0.7 Conclusion:

Objective: Identify which questions asked on Quora are duplicates of questions that have already been asked. This could be useful to instantly provide answers to questions that have already been answered. We are tasked with predicting whether a pair of questions are duplicates or not.

1. We are given a dataset with 4M rows and 5 columns: qid1, qid2, question1, question2, is_duplicate of which is_duplicate is a binary column containing whether the two questions are similar or not(1 or 0).
2. We perform EDA on the dataset to know about the distribution of the data and get more insights about the data.
3. Next we perform feature extraction to get some important features like length, common words which might be helpful to classify and then we go for advance feature extraction like fuzzy wuzzy.
4. We then preprocess the data to remove html tags etc and then visualize using wordclouds.

5. We then perform eda to know which of the features are actually important in classification using pairplots and histograms.
6. We convert the questions into vectors using tfidf glove vectorization and stack the features extracted to it (question1+question2+basic features+advance features).
7. We apply various model like logistic regression, linear svm and xgboost and compare them.
8. We use log loss and confusion matrix as the performance metric to compare various models.
9. we also try tfidf vectors on linear svm and logistic regression and hyperparameter tune xgboost model to furthur reduce the log loss.

```
In [1]: from prettytable import PrettyTable
```

```
x=PrettyTable()
```

```
x.field_names=['Algorithm','Vectorizer','alpha','max_depth','train log-loss','test log-loss']
x.add_row(["Logistic Regression","TFIDF-GLOVE",10,'-',0.513, 0.522 ])
x.add_row(["Linear SVM","TFIDF-GLOVE",0.0001,'-', 0.49,0.506])
x.add_row(["XGBOOST","TFIDF-GLOVE","-",4, 0.337,0.3518])
x.add_row(["Logistic Regression","TFIDF",0.0001 ,'-', 0.452,0.456])
x.add_row(["Linear SVM","TFIDF",0.00001,'-', 0.459,0.463])
x.add_row(["XGBOOST","TFIDF-GLOVE","-", 7,0.237,0.3348])
```

```
print(x)
```

Algorithm	Vectorizer	alpha	max_depth	train log-loss	test log-loss
Logistic Regression	TFIDF-GLOVE	10	-	0.513	0.522
Linear SVM	TFIDF-GLOVE	0.0001	-	0.49	0.506
XGBOOST	TFIDF-GLOVE	-	4	0.337	0.3518
Logistic Regression	TFIDF	0.0001	-	0.452	0.456
Linear SVM	TFIDF	1e-05	-	0.459	0.463
XGBOOST	TFIDF-GLOVE	-	7	0.237	0.3348

We can see that XgBoost with max depth of 7 and 100 estimators reduced the log loss to 0.334

```
In [ ]:
```